

None

Contents

Preface

[What Makes This Ebook Unique](#)

[Author Biography](#)

[Contacts and Social Media](#)

[Sources & References](#)

[Acknowledgments](#)

[Thank You!](#)

Contents(continued)

[Introduction](#)

[The "pivot_table" Syntax](#)

The "data" Parameter

The "values" Parameter

The "index" Parameter

The "columns" Parameter

The "aggfunc" Parameter

The "fill_value" Parameter

The "margins" Parameter

Contents(continued)

The "dropna" Parameter

The "margins_name" Parameter

The "observed" Parameter

The "sort" Parameter

Project(Real Life application)

Introduction

```
#Import Libraries
import pandas as pd, numpy as np
from IPython.display import HTML, display
```



Pivot Tables

With

Python *pandas*

sales_df

	Product	Price
0	CPU	2
1	CPU	1
2	Monitor	4
3	Monitor	3

```
pd.pivot_table(
    data = sales_df,
    index = 'Product',
    values = 'Price',
    aggfunc = 'sum'
)
```

	Price
Product	
CPU	3
Monitor	7



Kichere
The
Data Scientist

kicherethedatascientist@gmail.com

Mastering Pandas pivot tables offers a significant advantage over other tools in data analysis. Unlike other tools like Microsoft Excel, Google Sheets, Tableau, Power BI and QlikView/Qlik Sense and Looker, Pandas provides unparalleled flexibility and efficiency

when handling large datasets and complex data transformations. Its robust capabilities allow for seamless multi-indexing, advanced aggregations, and customized data summarization, empowering leaders and analysts to derive deeper insights swiftly and effectively.

Compared to alternatives, such as Excel or SQL, Pandas pivot tables excel in scalability and integration within Python's expansive data science ecosystem. This integration not only enhances workflow efficiency but also supports seamless transition between data manipulation, statistical analysis, and machine learning tasks, making it an indispensable tool for leaders striving to achieve data-driven decision-making.

To pivot data is to reorganize and summarize it in various ways, transforming detailed information into a more structured format.

	Category	Sales
0	A	100
1	A	150
2	B	200
3	B	250
4	C	300
5	C	350



Total Sales for each category:

	Sales
Category	
A	250
B	450
C	650

The "pivot_table" syntax

```
pandas.pivot_table(data, values=None, index=None, columns=None, aggfunc='mean',  
fill_value=None, margins=False, dropna=True, margins_name='All',  
observed=_NoDefault.no_default, sort=True) #
```

Parameters

```
"""  
Parameters:  
data  
values  
index  
columns  
aggfunc  
fill_value  
margins  
dropna  
margins_name  
observed  
sort  
"""
```

The **data** parameter

The "data" parameter accepts a DataFrame
Here "data" parameter is "sales_df"

Compute the total revenue generated from each product in the dataset (sales_df)

	Category	Sales
0	A	100
1	A	150
2	B	200
3	B	250
4	C	300
5	C	350

```
pd.pivot_table(  
    data=df,  
    index='Category',  
    values='Sales',  
    aggfunc='sum'  
)
```



Total Sales for each category:


	Sales
A	250
B	450
C	650

The **values** parameter

The "values" parameter accepts scalar values or list-like and is optional.

"values" as a scalar

values as a scalar, e.g. values = 'Price'

Product	Price		Product	Price
CPU	200		CPU	400
CPU	200			
Monitor	100		Monitor	200
Monitor	100			

Can you summarize the total sales value of each product in the dataset 'sales_df'?

values as a scalar:

A scalar is a single, indivisible value, such as a number or a string

sales_df

	Product	Price
0	CPU	200
1	CPU	200
2	Monitor	100
3	Monitor	100



Total sales summary:

	Price
Product	
CPU	400
Monitor	200

Code Snippet:

```
pd.pivot_table(
    data=sales_df,
    index='Product',
    values='Price',
    aggfunc='sum'
)
```


"values" as a list-like

"values" as a list-like e.g

List: [1, 2, 3] or List: ["Price", "Quantity"] for our example

Tuple: (1, 2, 3) or ("Price", "Quantity")

Set: {1, 2, 3}

String: "abc"

Dictionary: {'a': 1, 'b': 2}

NumPy Array: np.array([1, 2, 3])

Pandas Series: pd.Series([1, 2, 3])

Range: range(1, 4)

values=["Price", "Quantity"]

Display the total Quantity and Price for each Manager

values as a list-like, e.g. values=["Quantity", "Price"]

Manager	Product	Quantity	Price			Price	Quantity	
Debra Henley	CPU	2	600	→	Manager	Debra Henley	700	3
Debra Henley	Software	1	100					
Fred Anderson	CPU	1	300	→		Fred Anderson	600	4
Fred Anderson	Software	3	300					

Display the total Quantity and Price for each Manager

sales_df

	Manager	Product	Quantity	Price
0	Debra Henley	CPU	2	600
1	Debra Henley	Software	1	100
2	Fred Anderson	CPU	1	300
3	Fred Anderson	Software	3	300



Summary:

	Price	Quantity
Manager		
Debra Henley	700	3
Fred Anderson	600	4

Code Snippet:

```
pd.pivot_table(
    data=df,
    index="Manager",
    values=["Quantity", "Price"],
    aggfunc='sum'
)
```

#If you omit the values parameter, all columns that are not used as

#indexes or columns in the pivot table will be aggregated:

```
display(df)
pd.pivot_table(data = df,
               index= 'Product',
               #values = 'Price',
               aggfunc = 'sum'
              )
```

	Manager	Product	Quantity	Price
0	Debra Henley	CPU	2	600
1	Debra Henley	Software	1	100
2	Fred Anderson	CPU	1	300
3	Fred Anderson	Software	3	300

	Manager	Price	Quantity
Product			
CPU	Debra HenleyFred Anderson	900	3
Software	Debra HenleyFred Anderson	400	4

The **index** parameter

The "index" parameter accepts column name, Grouper, array, or lists of (column names, Groupers and arrays)

Compute the total sales revenue (Price) for each combination of manager and product in the DataFrame 'sales_df'?

sales_df

Pivot Table:

The **index** parameter:

Specify how data should be grouped

	Manager	Product	Quantity	Price
0	Debra Henley	CPU	2	600
1	Debra Henley	Software	1	100
2	Fred Anderson	CPU	1	300
3	Fred Anderson	Software	3	300



	Manager	Product	Price
	Debra Henley	CPU	600
		Software	100
	Fred Anderson	CPU	300
		Software	300

Therefore

	Manager	Product	Quantity	Price
0	Fred	RAM	2	600
1	Fred	CPU	1	100
2	Debra	CPU	3	900
3	Debra	RAM	3	300



Summary:

	Manager	Product	Price
	Debra	CPU	900
		RAM	300
	Fred	CPU	100
		RAM	600

Code snippet:

```
pd.pivot_table(
    sales_df,
    index=['Manager', 'Product'],
    values='Price',
    aggfunc={"Price": sum, "Product": len}
)
```

The **columns** parameter

The "columns" parameter accepts column name, Grouper, array, or lists of (column names, Groupers and arrays)

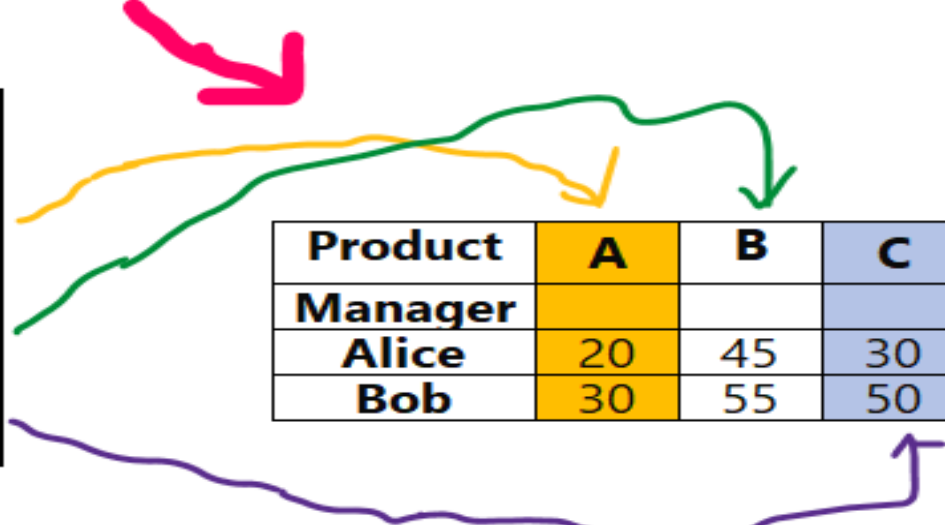
Create a pivot table to show the total Price for each Manager by Product in the DataFrame df?

The `columns` parameter

`pd.pivot_table(data = df, values='Price', index='Manager', columns='Product',)`

df

Manager	Product	Price
Alice	A	20
Bob	A	30
Alice	B	45
Bob	B	55
Alice	C	30
Bob	C	50



Product	A	B	C
Manager			
Alice	20	45	30
Bob	30	55	50

```
# Sample data
data = {
    'Manager': ['Alice', 'Bob', 'Alice', 'Bob', 'Alice', 'Bob'],
    'Product': ['A', 'A', 'B', 'B', 'C', 'C'],
    'Price': [20, 30, 45, 55, 30, 50]
}

df = pd.DataFrame(data)
df
```

	Manager	Product	Price
0	Alice	A	20
1	Bob	A	30
2	Alice	B	45
3	Bob	B	55
4	Alice	C	30

	Manager	Product	Price
5	Bob	C	50

```
pd.pivot_table(data=df, values='Price', index='Manager', columns='Product')
```

Product	A	B	C
Manager			
Alice	20.0	45.0	30.0
Bob	30.0	55.0	50.0

```
data = {
    'Manager': ['Debra Henley', 'Debra Henley', 'Debra Henley', 'Debra Henley', 'Debra Henley',
                'Debra Henley', 'Debra Henley', 'Debra Henley', 'Debra Henley', 'Fred Anderson',
                'Fred Anderson'],
    'Product': ['CPU', 'Software', 'Maintenance', 'CPU', 'CPU', 'CPU', 'Software', 'Maintenance',
                'CPU', 'CPU', 'CPU'],
    'Price': [30000, 10000, 5000, 35000, 65000, 40000, 10000, 5000, 35000, 65000, 30000],
    'Status': ['presented', 'presented', 'pending', 'declined', 'won', 'pending', 'presented',
                'pending', 'declined', 'won', 'presented']
}

df = pd.DataFrame(data);df
```

	Manager	Product	Price	Status
0	Debra Henley	CPU	30000	presented
1	Debra Henley	Software	10000	presented
2	Debra Henley	Maintenance	5000	pending
3	Debra Henley	CPU	35000	declined
4	Debra Henley	CPU	65000	won
5	Debra Henley	CPU	40000	pending
6	Debra Henley	Software	10000	presented
7	Debra Henley	Maintenance	5000	pending

	Manager	Product	Price	Status
8	Debra Henley	CPU	35000	declined
9	Fred Anderson	CPU	65000	won
10	Fred Anderson	CPU	30000	presented

```
pd.pivot_table(df, index='Manager', values='Price', columns=['Product', 'Status'], aggfunc='sum')
```

Product					CPU	Maintenance	Software
Status	declined	pending	presented	won		pending	presented
Manager							
Debra Henley	70000.0	40000.0	30000.0	65000.0		10000.0	20000.0
Fred Anderson	NaN	NaN	30000.0	65000.0		NaN	NaN

The **aggfunc** parameter

You can pass a function, a list of functions, a dictionary, or use the default "mean"

Compute average sales price per manager

When the `aggfunc` not provided, it calculates the "mean" by default.

sales_df

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300



Pivot Table:

	Price
Manager	
Debra	350.0
Fred	300.0

Code Snippet:

```
pd.pivot_table(
    data=sales_df,
    index="Manager",
    values='Price'
)
```

Can you summarize total sales by manager?

When you pass a function e.g "sum"

sales_df:

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300



Pivot Table:

	Price
Manager	
Debra	700
Fred	600

Code Snippet:

```
pd.pivot_table(
    sales_df,
    index='Manager',
    values='Price',
    aggfunc= 'sum'
)
```

Compute the highest and total values of 'Price' for each 'Manager'?

sales_df:

Pivot Table:

Code Snippet:

When you pass a
list of functions
e.g. [sum, max]

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300



		sum	max
	Price	Price	
Manager			
Debra		700	600
Fred		600	300

```
pd.pivot_table(
    sales_df,
    index='Manager',
    values='Price',
    aggfunc=['sum', 'max']
)
```

You could also find 'sum', 'max', 'min', 'mean', 'median', 'count', 'std', 'var'

calculates the total sales price and maximum product grouped by manager and product

When you pass a
dictionary e.g.
"Price": sum,
"Product": len

sales_df

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300



Pivot Table:

	Manager	Product	Price	Product
	Debra	CPU	600	CPU
		RAM	100	RAM
	Fred	CPU	300	CPU
		RAM	300	RAM

Pivot Table:

```
pd.pivot_table(
    sales_df,
    index=['Manager', 'Product'],
    values='Price',
    aggfunc={"Price": "sum", 'Prod
    })
```

You could also find 'sum', 'max', 'min', 'mean', 'median', 'count', 'std', 'var'

The **fill_value** parameter

The "fill_value" parameter specifies the value to replace missing values (NaN) in the resulting pivot table. It accepts a scalar value (single value), if no other value is provided, None is used as the default value.

*

Summarize the total 'Quantity' of products, grouped by 'Manager' across different 'Product' and 'Price' combinations

Dou you see the (NaNs) in a resulting pivot table?

sales_df:

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300



Pivot Table:

Product		CPU		RAM
Price	300	600	100	300
Manager				
Debra	NaN	2.0	1.0	NaN
Fred	1.0	NaN	NaN	3.0

**

Summarize the total 'Quantity' of products, grouped by 'Manager' across different 'Product' and 'Price' combinations

Replace the (NaNs) with a double dash (-)

sales_df:

Product	CPU		RAM	
Price	300	600	100	300
Manager				
Debra	NaN	2.0	1.0	NaN
Fred	1.0	NaN	NaN	3.0

Code Snippet:

```
pd.pivot_table(
    sales_df,
    index=["Manager"],
    columns=['Product', 'Price'],
    values='Quantity',
    aggfunc='sum',
    fill_value='--'
)
```



Pivot Table:

Product	CPU		RAM	
Price	300	600	100	300
Manager				
Debra	--	2	1	--
Fred	1	--	--	3

The **margins** parameter

**The "margins" parameter adds totals for each row and column.
It boolean and defaults to False, adds totals for each row and column when set to True.**

*

Show the total sales (sum of 'Price') for each manager across different products.

sales_df:

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300

**Without margins or
when `margins=False` .**

Code Snippet:

```
pd.pivot_table(
    sales_df,
    index='Manager',
    columns='Product',
    values='Price',
    aggfunc='sum'
)
```

Pivot Table:

Product	CPU	RAM
Manager		
Debra	600	100
Fred	300	300

*

Show the total sales (sum of 'Price') for each manager across different products

sales_df:

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300

**If `margins=True`,
special `All` columns and
rows will be added with
partial group aggregates
across the categories on
the rows and columns.**

Code Snippet:

```
pd.pivot_table(
    sales_df,
    index='Manager',
    columns='Product',
    values='Price',
    aggfunc='sum',
    margins=True
)
```

Pivot Table:

Product	CPU	RAM	All
Manager			
Debra	600	100	700
Fred	300	300	600
All	900	400	1300

The `dropna` parameter

The "dropna" parameter does not include columns whose entries are all NaN. If True, rows with a NaN value in any column will be omitted before computing margins. It is a boolean option and defaults to True.

*

Show the total sales (sum of 'Price') for each manager across different products.

Code Snippet:

If **margins=True** (default), Pandas excludes rows or columns with NaN values before computing a pivot table. This ensures NaNs do not affect calculations, and the resulting table is based only on available data without NaNs.

sales_df:

	Manager	Product	Quantity	Price	Status
0	Debra	CPU	2.0	600.0	None
1	Debra	RAM	NaN	100.0	None
2	Fred	CPU	1.0	300.0	None
3	Fred	RAM	3.0	NaN	None
4	None	None	NaN	NaN	None

```
pd.pivot_table(
    sales_df,
    index='Manager',
    columns='Product',
    values='Price',
    aggfunc='sum',
    margins=True,
    dropna=True
)
```

Pivot Table:

	Product	CPU	RAM	All
Manager				
Debra		600.0	100.0	700.0
Fred		300.0	0.0	300.0
All		900.0	100.0	1000.0

**

With **dropna=False**, NaN values are included in the pivot table

sales_df:

	Manager	Product	Quantity	Price	Status
0	Debra	CPU	2.0	600.0	None
1	Debra	RAM	NaN	100.0	None
2	Fred	CPU	1.0	300.0	None
3	Fred	RAM	3.0	NaN	None
4	None	None	NaN	NaN	None

Code Snippet:

```
pd.pivot_table(
    sales_df,
    index='Manager',
    columns=['Product', 'Status'],
    values='Price',
    aggfunc='sum',
    margins=True,
    dropna=False
)
```

Pivot Table:

	Product	CPU	RAM	NaN	All
Status		NaN	NaN	NaN	
Manager					
Debra		600.0	100.0	NaN	700.0
Fred		300.0	0.0	NaN	300.0
NaN		NaN	NaN	0.0	NaN
All		NaN	NaN	NaN	1000.0

The `margins_name` parameter

The "`margins_name`" parameter specifies the name of the row or column that will contain the totals when `margins=True`.

By default, the name is set to 'All'.

*

Show the total sales (sum of 'Price') for each manager across different products

The `margins_name` parameter:

Allows customization of the label used for the

totals row or column when `margins=True`.

This parameter is useful for providing clear and descriptive labels in pivot table summaries.

sales_df:

	Manager	Product	Quantity	Price
0	Debra	CPU	2	600
1	Debra	RAM	1	100
2	Fred	CPU	1	300
3	Fred	RAM	3	300

Code Snippet:

```
pd.pivot_table(
    sales_df,
    index='Manager',
    columns='Product',
    values='Price',
    aggfunc='sum',
    margins=True,
    margins_name='Total'
)
```

Pivot Table:

Product	CPI
Manager	
Debra	600
Fred	300
Total	900

The `sort` parameter

The "`sort`" parameter specifies if the result should be sorted

By default, "`sort`" is set to True

Price of products sold, grouped by manager and product?

sales_df:

Code Snippet:

Pivot Table:

Setting (sort=True)

(default), ensures that the data is presented in an organized manner based on the index labels 'Manager' and 'Product'.

	Manager	Product	Quantity	Price
0	Fred	RAM	2	600
1	Fred	CPU	1	100
2	Debra	CPU	3	900
3	Debra	RAM	3	300

```
pd.pivot_table(
    sales_df,
    values='Price',
    index=['Manager', 'Product'],
    aggfunc='sum',
    sort=True
)
```

		Price
Manager	Product	
Debra	CPU	900
	RAM	300
Fred	CPU	100
	RAM	600

Price of products sold, grouped by manager and product?

Setting**(sort=False)**

maintains the original order of the data

sales_df:

	Manager	Product	Quantity	Price
0	Fred	RAM	2	600
1	Fred	CPU	1	100
2	Debra	CPU	3	900
3	Debra	RAM	3	300

Code Snippet:

```
pd.pivot_table(
    sales_df,
    values='Price',
    index=['Manager', 'Product'],
    aggfunc='sum',
    sort=False
)
```

Pivot Table:

		Price
Manager	Product	
Debra	CPU	600
	RAM	100
Fred	CPU	300
	RAM	300

The **observed** parameter

**The "observed" parameter is Deprecated since version 2.2.0:
my pandas version = 2.2.2**

```
#print(pd.__version__)
#2.2.2
```

```
data = {
    'Region': ['East', 'East', 'West', 'West', 'North'],
    'Salesperson': ['Alice', 'Bob', 'Alice', 'Charlie', 'David'],
    'Sales': [10000, 15000, 12000, 8000, 9000]
}

df = pd.DataFrame(data)
df
```

	Region	Salesperson	Sales
0	East	Alice	10000
1	East	Bob	15000
2	West	Alice	12000
3	West	Charlie	8000
4	North	David	9000

```
# Pivot table with observed=True
df.pivot_table(index='Region', columns='Salesperson',
                values='Sales', aggfunc='sum', observed=True)
```

Salesperson	Alice	Bob	Charlie	David
Region				
East	10000.0	15000.0	NaN	NaN
North	NaN	NaN	NaN	9000.0
West	12000.0	NaN	8000.0	NaN

```
# Pivot table with observed=False
df.pivot_table(index='Region', columns='Salesperson',
                values='Sales', aggfunc='sum', observed=False)
```

Salesperson	Alice	Bob	Charlie	David
Region				
East	10000.0	15000.0	NaN	NaN
North	NaN	NaN	NaN	9000.0
West	12000.0	NaN	8000.0	NaN

#Deprecated since version 2.2.0: The default value of False is deprecated and #will change to True in a future version of pandas.

Contacts and Social Media

Kichere Magubu

 **Dar es salaam, Tanzania**

 **Kichere Magubu**

 **Kichere The Data Scientist**

 **KichereTheDataScientist**

 **KichereTheDataScientist**

 **kicherethedatascientist@gmail.com**

 **kicherethedatascientist**

 **+255 654 729 851**

Powered by
Eastern Africa Statistical Training Centre
SKT Tanzania Ltd

Why This E-book?

"The aim of this ebook is to give you the 'aha' moment right away concept."

- **Practical step By Step Guide With Simple Examples**
- **Visual Illustrations and Interactive**
- **Simple Datasets**
- **Comprehensive Coverage**
- **Designed for Beginners**
- **No Prior Knowledge Required**
- **Includes Pandas Documentation**

Project(Real Life application)

Practical Business Python

[Pandas Pivot Table Explained](https://pbpython.com/pandas-pivot-table-explained.html) (<https://pbpython.com/pandas-pivot-table-explained.html>)

Acknowledgments

First and foremost, I express my gratitude to God for His blessings and guidance throughout this journey. I am also deeply thankful to the Eastern Africa Statistical Training Centre and SKT Tanzania Ltd for their invaluable support. Special thanks to my parents whose unwavering encouragement and support have been the cornerstone of this endeavor.

Sources & References

[pandas.pivot_table Documentation](https://pandas.pydata.org/docs/reference/api/pandas.pivot_table.html)

(https://pandas.pydata.org/docs/reference/api/pandas.pivot_table.html)

[Pandas Pivot Table Explained](https://pbpython.com/pandas-pivot-table-explained.html) (<https://pbpython.com/pandas-pivot-table-explained.html>)

Author Biography

Kichere Magubu is a data enthusiast and content creator.

Thank you!

Thank you for reading this e-book!.

It is designed to assist you in understanding pivot tables using Python pandas. If you found this book valuable and beneficial, I kindly request that you take a moment to review it. Your honest feedback is highly appreciated and makes a significant difference to me.

Please take a minute to write me an honest review. Your support means the world to me!

Thank you!

```
#print("The cell to convert jupyter notebook to html")  
!jupyter nbconvert --to hide_code_html "Pivot Tables.ipynb"
```